

AI-POWERED LEGAL DOCUMENT SUMMARIZATION AND CASE PRECEDANT FINDER

PHASE 1 REPORT

Submitted by

SANJAY KISHAN D 230701287

SANJEEV ADTIYA P 230701292

SANJAY R 230701291

SANJAY G 230701286

In partial fulfillment for the award of the degree

of

BACHELOR OF ENGINEERING

IN

COMPUTER SCIENCE AND ENGINEERING



RAJALAKSHMI ENGINEERING COLLEGE

ANNA UNIVERSITY: CHENNAI 600025

OCTOBER 2026

ANNA UNIVERSITY, CHENNAI

BONAFIDE CERTIFICATE

Certified that this Report titled "**AI-POWERED LEGAL DOCUMENT SUMMARIZATION AND CASE PRECEDANT FINDER**" is the bonafide work of **Sanjay Kishan D (230701287), Sanjeev Adtiya P (230701292), Sanjay R (230701291), Sanjay G (230701286)** who carried out the work under my supervision. Certified further that to the best of my knowledge the work reported herein does not form part of any other thesis or dissertation on the basis of which a degree or award was conferred on an earlier occasion on this or any other candidates.

This project addresses the fallowing Sustainable Development Goal(S):
Goal No :16 - Peace, Justice and Strong Institutions

SIGNATURE

Dr. J. Manoranjini. Ph.D.,

Professor and Head,

Department of Computer Science and

Engineering,

Rajalakshmi Engineering College,

Chennai-602 105.

SIGNATURE

Ms. R. Akiladevi, M.E.,

Supervisor,

Assistant Professor (SS),

Department of Computer Science and

Engineering,

Rajalakshmi Engineering College,

Chennai-602 105.

Submitted to Project Viva-Voce Examination held on _____

Internal Examiner

External Examiner

ABSTRACT

Legal professionals and individuals often face significant challenges when reviewing lengthy legal documents, contracts, statutes, and court cases. Identifying important clauses, understanding complex legal content, and finding relevant case precedents can be time-consuming and require extensive manual effort. The proposed project, AI-Powered Legal Document Summarization and Case Precedent Finder, introduces an AI-powered legal document analysis framework that combines legal text processing, domain-specific language understanding, semantic precedent retrieval, and evidence-grounded summarization. The system accepts legal documents in PDF and DOCX formats and uses Legal-BERT to identify important clauses, entities, rights, obligations, and liability-related information. A semantic retrieval module uses Sentence-Transformers to convert legal content into meaningful vector representations and applies similarity-based search to identify and rank relevant court precedents. The retrieved case laws and statutes are then used as supporting evidence for generating simplified legal summaries and insights. A Streamlit-based interactive dashboard presents the extracted findings, retrieved precedents, and generated outputs in a user-friendly manner. Unlike traditional legal database systems that primarily depend on manual searching and document analysis, the proposed framework integrates domain-specific NLP, semantic search, and retrieval-augmented generation to streamline legal research and document understanding. The system also incorporates grounded-versus-ungrounded evaluation to identify unsupported information and fabricated precedents. By combining Legal-BERT, semantic precedent retrieval, and evidence-grounded generation, the proposed system aims to reduce manual effort, improve access to relevant legal information, and provide more reliable and understandable analysis of complex legal documents.

Keywords: Artificial Intelligence, Legal-BERT, Natural Language Processing (NLP), Legal Document Summarization, Legal Document Analysis, Semantic Search.

ACKNOWLEDGEMENT

Initially we thank the Almighty for being with us through every walk of our life and showering his blessings through the endeavor to put forth this report. Our sincere thanks to our Chairman **Mr. S. MEGANATHAN, B.E, F.I.E.**, our Vice Chairman **Mr. ABHAY SHANKAR MEGANATHAN, B.E., M.S.**, and our respected Chairperson **Dr. (Mrs.) THANGAM MEGANATHAN, Ph.D.**, for providing us with the requisite infrastructure and sincere endeavoring in educating us in their premier institution.

Our sincere thanks to **Dr. S.N. MURUGESAN, M.E., Ph.D.**, our beloved Principal for his kind support and facilities provided to complete our work in time. We express our sincere thanks to **Dr. J. Manoranjini Ph.D**, Professor and Head of the Department of Computer Science and Engineering for her guidance and encouragement throughout the project work. We convey our sincere and deepest gratitude to our internal guide, **Ms. R. Akiladevi, M.E.**, Assistant Professor (SS), Department of Computer Science and Engineering for her valuable guidance throughout the course of the project. We thank our Project Coordinator, **DR. S. MADHUSUDHANAN.**, Associate Professor, Department of Computer Science and Engineering for his useful tips during our review to build our project.

SANJEEV ADTHIYA: 230701292

SANJAY KISHAN D: 230701287

SANJAY G: 230701286

SANJAY R: 230701291

TABLE OF CONTENT

CHAPTER NO.	TITLE	PAGE NO..
	ABSTRACT	iv
	ACKNOWLEDGEMENT	v
	LIST OF FIGURES	viii
	LIST OF ABBREVIATIONS	ix
1.	INTRODUCTION	10
	1.1 GENERAL	10
	1.2 OBJECTIVE	11
	1.3 EXISTING SYSTEM	12
	1.4 PROPOSED SYSTEM	13
2.	LITERATURE SURVEY	15
	2.1 LITERATURE REVIEW	15
	2.2 COMPARISON AND DISCUSSION	15
	2.3 CONCLUSION	15
3.	SYSTEM DESIGN	21
	3.1 SYSTEM ARCHITECTURE	23
	3.2 SYSTEM REQUIREMENTS (3 at s)	24
	3.2.1 SOFTWARE REQUIREMENTS	25
	3.2.2 HARDWARE REQUIREMENTS	26
	3.3 SYSTEM DESIGN	27
	3.3.1 ACTIVITY DIAGRAM	30
	3.3.2 DATA FLOW DIAGRAM	
4.	PROJECT DESCRIPTION	32
	4.1 METHODOLOGIES	32
	4.2 MODULES	33
	4.2.1 DOCUMENT PROCESSING MODULE	33

	4.2.1 LEGAL-BERT INFORMATION EXTRACTION MODULE	33
	4.2.2 SEMANTIC PRECEDENT RETRIEVAL MODULE REINFORCEMENT PACKAGE	34
	4.2.3 EVIDENCE-GROUNDED SUMMARIZATION MODULE	34
	4.2.4 INTERACTIVE DASHBOARD MODULE	35
5.	IMPLEMENTATION AND RESULT DISCUSSION	36
	5.1 IMPLEMENTATION RESULTS	37
	5.1.1 IMPLEMENTATION OF DATA PRE-PROCESSING	
	5.1.2 IMPLEMENTATION OF DATA TOKENIZATION	39
	5.1.3 IMPLEMENTATION OF DATA EMBEDDING	
	5.2 EXPECTED RESULT	40
6.	CONCLUSION AND FUTURE WORK	46
	6.1 CONCLUSION	46
	6.2 FUTURE WORK	47
	REFERENCES	

LIST OF FIGURES

Figure No.	Title	Page No.
3.1	System Architecture of the AI-Powered Legal Document Summarization and Case Precedent Finder	22
3.2	End-to-End System Workflow of the Legal Document Analysis and Precedent Retrieval System	29
3.3	Data Flow Diagram of the AI-Powered Legal Document Summarization and Case Precedent Finder	31
5.1	Implementation of UI Design	37
5.2	Text Input for case ID Retrieval	38
5.3	Legal Document Data Pre-processing Pipeline	39
5.4	Legal Text Embedding and Feature Vectorization Architecture	41
5.5	Wait-Time Prediction Error (MAE in Minutes) Chart	43
Table 5.1	Performance Evaluation Matrix	42

LIST OF ABBREVIATION

Abbreviation	Description
AI	Artificial Intelligence (CNETRALIZE)
CNN	Convolutional Neural Network
DFD	Data Flow Diagram
ETA	Expected Time of Arrival
GPS	Global Positioning System
IoT	Internet of Things
LSTM	Long Short-Term Memory
MAE	Mean Absolute Error
ML	Machine Learning
MSE	Mean Square Error
NLP	Natural Language Processing
ReLU	Rectified Linear Unit
RMSE	Root Mean Square Error
RNN	Recurrent Neural Network
UI	User Interface
UML	Unified Modeling Language
API	Application Programming Interface
CNN-LSTM	Hybrid Deep Learning Architecture combining CNN and LSTM
DNN	Deep Neural Network
STN	Smart Transportation Network
TPU	Tensor Processing Unit
GPU	Graphics Processing Unit

CHAPTER 1

INTRODUCTION

1.1 GENERAL

Legal professionals, researchers, and individuals often need to review large volumes of legal documents, contracts, statutes, regulations, and court cases to understand relevant legal information. However, conventional legal research and document analysis largely depend on manual reading and database searching, which can be time-consuming, especially when dealing with lengthy and complex legal documents. Identifying important clauses, extracting relevant entities and obligations, understanding legal content, and finding related case precedents can require considerable effort. Traditional legal search systems may also require users to manually examine numerous results before identifying the most relevant precedents.

The proposed AI-Powered Legal Document Summarization and Case Precedent Finder aims to address these challenges by integrating Artificial Intelligence and Natural Language Processing with legal document analysis. The system accepts legal documents in PDF and DOCX formats and uses Legal-BERT to identify important clauses, entities, rights, obligations, and other relevant legal information. A semantic precedent retrieval module uses Sentence-Transformers to convert legal content into meaningful vector representations and applies similarity-based retrieval to identify and rank relevant court precedents.

The retrieved legal evidence is then used to support the generation of simplified summaries and legal insights. The system emphasizes evidence-grounded responses, where relevant statutes and court judgments are used to support generated information and reduce unsupported or fabricated legal claims. A Streamlit-based interactive dashboard provides users with extracted findings, retrieved precedents, and generated summaries in an accessible format. The

overall workflow consists of

Legal Content → Clause and Entity Extraction → Semantic Precedent Retrieval
→ Evidence-Grounded Summary → Dashboard.

By combining domain-specific legal language processing, semantic search, and evidence-grounded generation, the proposed system aims to reduce the manual effort involved in legal document analysis and precedent discovery. It provides a structured approach for understanding complex legal documents and accessing relevant case information, while incorporating evaluation of grounded and ungrounded responses to identify unsupported information.

1.2 OBJECTIVE

The main goal of the AI-Powered Legal Document Summarization and Case Precedent Finder is to design and develop an intelligent legal document analysis framework that leverages Artificial Intelligence and Natural Language Processing to simplify legal document understanding and improve the discovery of relevant case precedents. The objectives are outlined as follows:

Legal Document Processing: To develop a document processing module capable of accepting legal documents in PDF and DOCX formats and converting complex legal content into structured information for further analysis.

1. Legal Information Extraction: To utilize Legal-BERT for identifying and extracting important clauses, entities, rights, obligations, and liability-related information from legal documents.
2. Semantic Case Precedent Retrieval: To implement a semantic retrieval mechanism using Sentence-Transformers and vector similarity to identify and rank relevant court precedents based on the meaning and context of the legal document.
3. Evidence-Grounded Legal Summarization: To generate simplified legal summaries and insights using retrieved case laws and statutes as supporting evidence, thereby reducing unsupported or fabricated legal information.
4. Hallucination Evaluation: To evaluate the reliability of generated legal responses by comparing grounded and ungrounded outputs and identifying fabricated precedents or unsupported statutes.

5. Interactive Legal Analysis Dashboard: To provide an interactive Streamlit-based dashboard through which users can view extracted legal findings, retrieved precedents, and AI-generated summaries in an organized and user-friendly manner.

Achieving these objectives will transform conventional manual legal document analysis into a more structured and AI-assisted workflow that reduces the effort required for reviewing legal content, improves access to relevant precedents, and supports more reliable legal information retrieval and summarization.

1.3 EXISTING SYSTEM

Existing legal research and document analysis systems primarily rely on conventional legal databases, manual document review, and keyword-based searching to handle contracts, statutes, regulations, and court cases. Platforms such as Manupatra, Westlaw, and SCC Online provide access to extensive collections of legal information; however, users often need to manually search, review, and interpret the retrieved documents to identify the information relevant to their requirements.

Legal professionals also spend considerable time manually identifying important clauses, obligations, entities, and other relevant information from lengthy legal documents. As the volume and complexity of legal content increases, manual analysis becomes time-consuming and makes it difficult to efficiently locate the most relevant legal information. Traditional systems generally provide search results but offer limited assistance in automatically understanding the context and meaning of complex legal documents.

Another limitation of existing systems is the use of generic Large Language Models (LLMs) for legal summarization without sufficient evidence verification. Although generic LLMs can simplify complex legal content, they may generate unsupported claims or fabricated case laws when the generated information is not properly verified against authentic legal sources.

Furthermore, conventional legal systems generally lack an integrated workflow that combines domain-specific legal text extraction, semantic precedent retrieval, and evidence-grounded summarization. Users may therefore need to perform document analysis, precedent searching, and interpretation as separate activities. These limitations highlight the need for an AI-assisted legal analysis system that can extract relevant

information from legal documents, identify semantically related precedents, and support generated summaries with authentic legal evidence.

1.4 PROPOSED SYSTEM

The AI-Powered Legal Document Summarization and Case Precedent Finder introduces an intelligent and integrated legal analysis framework designed to overcome the limitations of conventional legal research and document analysis systems. The proposed system combines Legal Document Processing, Legal-BERT-based Information Extraction, Semantic Case Precedent Retrieval, Evidence-Grounded Summarization, and an Interactive Dashboard into a unified workflow. The document processing module accepts legal documents in PDF and DOCX formats and converts the legal content into structured information for further analysis.

The Legal-BERT Information Extraction Module analyzes the legal content to identify important clauses, entities, rights, obligations, and liability-related information. Legal-BERT is used as a domain-specific language model for understanding the characteristics of legal text. The extracted information is then passed to the Semantic Precedent Retrieval Module, where Sentence-Transformers convert the relevant legal content into vector representations. Cosine similarity is used to identify and rank court precedents that are semantically related to the input document.

The proposed system further incorporates Legal Evidence Verification and Evidence-Grounded Summarization. Retrieved case laws and statutes are used as supporting evidence for generated summaries and legal insights, helping reduce unsupported information and fabricated legal claims. The system also includes a hallucination evaluation mechanism that compares grounded and ungrounded responses to identify fabricated precedents and unsupported statutes.

A Streamlit-based interactive dashboard provides users with the extracted legal findings, retrieved precedents, and generated summaries in an organized interface. The overall workflow follows Legal Content → Clause and Entity Extraction → Semantic Precedent Retrieval → Evidence-Grounded Summary → Streamlit Dashboard.

The modular architecture allows the document parser, Legal-BERT extraction, retrieval layer, and generation components to work together as an integrated system. By combining domain-specific NLP, semantic search, and evidence-grounded generation,

the proposed system aims to reduce manual legal research effort, simplify complex legal documents, improve precedent discovery, and provide more reliable legal information. Overall, the system provides an AI-assisted and scalable approach to modern legal document analysis and case precedent research.

CHAPTER 2 LITERATURE SURVEY

The advancement of Artificial Intelligence, Natural Language Processing, and transformer-based language models has created new possibilities for intelligent legal document analysis and automated legal research. Traditional legal systems mainly depend on manual document review and database searching, which can make the identification of relevant clauses and case precedents time-consuming. Recent research has therefore focused on domain-specific language models, legal text understanding, semantic retrieval, and Retrieval-Augmented Generation (RAG) to improve the efficiency and reliability of legal information processing.

2.1. LITERATURE REVIEW

2.1.1 Legal-BERT for Legal Language Understanding

Chalkidis *et al.* introduced LEGAL-BERT, a domain-specific language model designed specifically for legal text processing. The model was developed to address the limitations of general-purpose language models when handling specialized legal terminology and concepts. Their work demonstrates the effectiveness of domain-specific pre-training for legal language understanding and provides an important foundation for applying transformer-based models to legal documents [7]. This work is directly relevant to the proposed system, where Legal-BERT is used to analyze legal documents and identify important legal information such as clauses, entities, rights, obligations, and liability-related content.

2.1.2 BERT for Natural Language Understanding

Devlin *et al.* introduced BERT (Bidirectional Encoder Representations from Transformers), a transformer-based language model capable of learning contextual representations by considering both left and right contexts of words [8]. BERT achieved strong performance across several natural language

processing tasks and established a foundation for subsequent domain-specific models. Legal-BERT extends this concept to the legal domain, making the work particularly relevant to the proposed legal document analysis system.

2.1.3 LexGLUE for Legal Language Understanding

Chalkidis *et al.* introduced LexGLUE, a benchmark designed to evaluate language models on a diverse set of legal language understanding tasks [5]. The benchmark demonstrates the complexity and diversity of legal NLP tasks and highlights the need for models capable of handling different types of legal text. This work supports the use of specialized transformer-based approaches in the proposed system for extracting meaningful information from legal documents.

2.1.4 Sentence-BERT for Semantic Similarity

Reimers and Gurevych proposed Sentence-BERT (SBERT), which generates efficient sentence-level embeddings that can be compared using similarity measures such as cosine similarity [16]. SBERT significantly improves the efficiency of semantic similarity and information retrieval tasks compared with directly comparing sentences using conventional BERT architectures. This approach is relevant to the proposed system because Sentence-Transformers are used to generate vector representations of legal text and retrieve semantically relevant case precedents.

2.1.5 Legal Document Retrieval

Shao *et al.* investigated legal document retrieval using deep neural representation learning, demonstrating the potential of neural representations for identifying relevant legal documents [17]. Traditional keyword-based retrieval may fail when legally related documents use different terminology or expressions. Representation-based retrieval can instead capture semantic relationships between legal documents. This concept supports the proposed semantic precedent retrieval module, which identifies relevant case laws based on the contextual meaning of legal content.

2.1.6 Legal Text Classification

Guo *et al.* studied the application of pre-trained language models for legal text classification [9]. Their work demonstrates the usefulness of transformer-based representations for categorizing legal documents and understanding domain-specific legal information. The findings are relevant to the proposed system because Legal-BERT is used as the foundation for processing and extracting meaningful information from legal documents.

2.1.7 Legal Text Summarization

Bhattacharya *et al.* investigated the use of deep learning techniques for legal text summarization [2]. Legal documents are often lengthy and contain complex terminology, making manual summarization time-consuming. Deep learning-based summarization can help reduce the amount of information that users need to manually review while preserving important legal content. This work supports the summarization component of the proposed system, which aims to generate concise and understandable summaries of lengthy legal documents.

2.1.8 Long-Document Processing

Beltagy *et al.* proposed Longformer, a transformer architecture designed specifically for processing long documents efficiently [1]. Conventional transformer architectures can face computational limitations when processing lengthy documents because of their attention mechanism. Long-document processing is particularly important in legal applications because contracts, statutes, judgments, and other legal documents can contain large amounts of text. This research provides useful background for designing systems capable of processing lengthy legal documents.

2.1.9 Retrieval-Augmented Generation

Guu *et al.* introduced retrieval-augmented language modeling, demonstrating how external information retrieval can be combined with language models to improve knowledge-intensive tasks [10]. Retrieval-based approaches

allow models to access relevant external information instead of depending entirely on knowledge stored within model parameters. This concept is relevant to the proposed system because retrieved case precedents and statutes are provided as supporting evidence before generating the final legal summary.

2.1.10 Human Verification and Legal AI

Research on legal AI and computational legal research emphasizes the importance of combining automated analysis with appropriate human interpretation [4]. Legal decisions and interpretations can have significant consequences, and therefore AI-generated outputs should be treated as decision-support information rather than as a replacement for legal professionals. The proposed system follows this principle by presenting extracted information, relevant precedents, and evidence-supported.

2.2 COMPARISON AND DISCUSSION

The reviewed literature demonstrates the increasing application of transformer-based models, semantic embeddings, information retrieval, and deep learning techniques in legal NLP. General-purpose models such as BERT provide strong contextual language representations [8], while Legal-BERT demonstrates the advantages of adapting transformer models to the legal domain [7].

For semantic retrieval, Sentence-BERT provides an efficient approach for generating sentence embeddings and calculating semantic similarity [16]. Legal document retrieval research further demonstrates the potential of neural representations for identifying relevant legal information [17]. These approaches provide the foundation for the proposed Case Precedent Retrieval Module.

Research on legal summarization demonstrates the potential of deep learning for reducing the complexity of lengthy legal documents [2]. Longformer further addresses the challenge of processing long documents efficiently [1]. However, summarization alone does not ensure that generated information is supported by reliable legal sources.

The proposed system addresses this limitation by combining Legal-BERT-based information extraction, semantic precedent retrieval, and evidence-grounded summarization. Relevant case precedents and legal sources are retrieved before the generation stage, allowing the generated summary to be supported by external legal evidence.

Overall, the literature indicates that existing research provides strong solutions for individual legal NLP tasks such as classification, semantic similarity, retrieval, and summarization. However, there remains an opportunity to integrate these capabilities into a single workflow. The proposed system attempts to bridge this gap by combining domain-specific legal understanding with semantic precedent retrieval and evidence-grounded legal summarization.

2.3. CONCLUSION

In summary, the reviewed studies demonstrate the growing role of Artificial Intelligence, Natural Language Processing, and transformer-based models in modern legal document analysis and research. Domain-specific models such as Legal-BERT can support the understanding and extraction of information from legal texts, while semantic retrieval techniques can assist in identifying relevant case precedents. Research on Retrieval-Augmented Generation (RAG) further demonstrates the potential of retrieved evidence to improve the reliability of generated outputs and reduce unsupported information.

The proposed AI-Powered Legal Document Summarization and Case Precedent Finder aims to bring these capabilities together by combining legal document processing, Legal-BERT-based information extraction, semantic precedent retrieval, evidence verification, and AI-assisted summarization within a unified framework. The system is designed to help users identify important legal information, discover relevant court precedents, and understand complex legal documents through simplified, evidence-supported outputs.

By integrating these capabilities with an interactive Streamlit dashboard, the proposed system provides a structured and AI-assisted approach to legal research

and document analysis. The emphasis on evidence-grounded responses and hallucination evaluation further supports the goal of producing outputs that are connected to authentic legal sources. Overall, the proposed framework aims to make legal document analysis more efficient, accessible, structured, and evidence-oriented, while assisting users in understanding legal content and discovering relevant precedents.

CHAPTER 3

SYSTEM DESIGN

System design is the process of translating the requirements of the AI-Powered Legal Document Summarization and Case Precedent Finder into a structured technical framework. It defines the system architecture, major modules, processing workflow, and flow of information between different components. The proposed system integrates Legal Document Processing, Legal-BERT-based Information Extraction, Semantic Case Precedent Retrieval, Legal Evidence Verification, and Evidence-Grounded Summarization. It also provides an interactive Streamlit dashboard for presenting extracted findings, retrieved precedents, and generated summaries. The design focuses on efficient legal information processing, semantic understanding, evidence-based generation, reliability, scalability, and ease of use.

3.1 SYSTEM ARCHITECTURE

The AI-Powered Legal Document Summarization and Case Precedent Finder is designed as a modular legal analysis framework that processes legal documents through a sequence of specialized components. The system accepts legal text and documents in PDF and DOCX formats and processes them through a document parsing module. The parsed content is then passed to the Legal-BERT module, which extracts important clauses, entities, rights, obligations, and liability-related information. These extracted findings form the basis for subsequent legal analysis and precedent retrieval.

The Document Processing and Legal Information Extraction Layer forms the initial processing stage of the system. The `doc_parser.py` module is responsible for converting uploaded legal documents into structured content, while `legal_bert.py` performs legal-domain text analysis and information extraction. This layer transforms unstructured legal documents into meaningful legal findings that can be used by the downstream retrieval and generation components.

The Semantic Precedent Retrieval Layer identifies relevant court cases based on the contextual meaning of the extracted legal information. Sentence-Transformers are used to convert legal content into vector representations, and similarity-based retrieval is applied to find and rank relevant court precedents. The retrieval.py module performs this retrieval process, allowing the system to identify cases that are semantically related to the input legal document rather than relying only on direct keyword matching.




Fig.3.1 System Architecture of the AI-Powered Legal Document Summarization and Case Precedent Finder

The LLM Synthesis and Evidence Verification Layer uses retrieved case laws and statutes as supporting evidence for generating legal summaries and insights. The system is designed to ground generated responses in authentic legal evidence to reduce unsupported information and fabricated legal claims. The synthesis component uses Gemini, while the system also evaluates grounded and ungrounded responses to identify unsupported statutes and fabricated precedents. The Presentation Layer consists of a Streamlit-based interactive dashboard through which users can access the processed legal information. The dashboard

presents extracted findings, retrieved case precedents, and AI-generated summaries in an organized manner. The overall system architecture follows the workflow:

Legal Text/Document → Document Parser → Legal-BERT → Semantic Retrieval → Gemini Synthesis → Report/Dashboard.

The modular architecture allows individual components such as document parsing, legal information extraction, precedent retrieval, and response generation to perform their respective functions while contributing to a unified legal analysis workflow. The architecture therefore provides a structured foundation for reducing manual legal research effort and supporting efficient analysis of complex legal documents.

3.2 SYSTEM REQUIREMENTS

The AI-Powered Legal Document Summarization and Case Precedent Finder is designed to process legal documents, extract relevant legal information, retrieve related case precedents, and generate evidence-grounded summaries through an interactive dashboard. To ensure proper development, deployment, and operation, the system requirements can be classified into functional and non-functional requirements.

Hardware and Software Requirements

To ensure smooth development, deployment, and operation of the AI-Powered Legal Analyzer, the following hardware and software specifications are recommended:

3.2.1 Hardware Requirements

1. Server/Processing Unit

- Multi-core processor (Intel i5/i7 or equivalent) for handling backend services and AI-based processing. Minimum 16 GB RAM for efficient execution of machine learning algorithms.
- Minimum 16 GB RAM for efficient execution of AI models, databases, and multiple services.
- SSD storage (500 GB or more) for storing patient records, consultation data, pharmacy information, and AI models.

2. User Devices

- Smartphones or tablets for patients and doctors to access their respective mobile applications.
- Desktop/laptop computers for clinic assistants to manage appointments, queues, patient records, inventory, and reports.

3. Networking

- Stable internet connectivity for communication between frontend applications, backend services, and external services.
- Secure network connectivity to support HTTPS REST APIs and Socket.IO for real-time communication and queue updates.

4. Supporting Devices

- a. Adequate local or cloud-based storage should be available for maintaining PDF/DOCX legal documents, extracted findings, retrieved case information, and generated reports.
- b. Storage should also support legal model files, embeddings/vector data, and other artifacts required by the retrieval and analysis pipeline.

The hardware setup enables the proposed system to efficiently process legal documents, perform Legal-BERT-based information extraction, semantic precedent retrieval, and AI-assisted summarization. Adequate processing power, memory, storage, and network connectivity allow the system to handle the different stages of the legal analysis workflow while providing flexibility for future expansion. The system architecture presented in the project follows the flow from legal document input through parsing, Legal-BERT extraction, retrieval, Gemini synthesis, and final dashboard/report output.

3.2.2 Software Requirements

1. Operating System

- Windows 10/11 or Linux (Ubuntu 20.04 or later) for development, testing, and server environments.

2. Development Tools and Languages

- Python 3.x for implementing AI-based patient triage, prediction, and data processing modules.
- JavaScript/TypeScript for developing frontend applications and interactive web interfaces.
- Node.js with Express.js for implementing backend services and RESTful APIs.
- React.js for developing responsive web applications and clinic management portals.

3. Libraries and Frameworks

- Machine Learning: Scikit-learn and related Python libraries for symptom classification, recommendations, and predictive analysis.
- Data Processing: Pandas and NumPy for data preprocessing, analysis, and manipulation.
- Backend Communication: Socket.IO for real-time queue updates and notifications.

4. Database Management System

- PostgreSQL/MySQL for storing patient details, doctor information, appointments, consultations, prescriptions, and inventory records.
- File/Cloud Storage for maintaining consultation documents, prescriptions, reports, and other supporting files.

5. Additional Tools

- Version Control: Git/GitHub for source-code management, collaboration, and version tracking.
- API Documentation: Swagger/OpenAPI for documenting and testing RESTful backend services.
- External APIs: Maps and communication APIs for location services, distance estimation, email, SMS, and notification delivery.

The recommended software stack enables the AI-Powered Legal Document Summarization and Case Precedent Finder to efficiently support legal document processing, semantic information retrieval, and evidence-grounded summarization. By combining Legal-BERT, document processing modules, Sentence-Transformers, semantic retrieval, Gemini, and the Streamlit dashboard, the system can process

legal documents, extract relevant clauses and entities, identify related case precedents, and present the resulting analysis in an organized manner.

Furthermore, this configuration provides flexibility for integrating additional legal datasets, AI models, retrieval techniques, and evaluation mechanisms in the future. The modular architecture allows individual components such as document parsing, Legal-BERT extraction, precedent retrieval, and LLM synthesis to be enhanced independently, supporting the scalability, maintainability, and adaptability of the proposed legal analysis system.

3.3. SYSTEM DESIGN

The AI-Powered Legal Document Summarization and Case Precedent Finder design model describes how legal documents are transformed from raw input into structured legal findings, relevant case precedents, and evidence-grounded summaries through a sequence of processing stages. The system accepts legal text and documents in PDF and DOCX formats and performs document parsing and preprocessing before passing the processed content to the Legal-BERT Information Extraction Module. Legal-BERT identifies important clauses, entities, rights, obligations, and liability-related information. The extracted information is then provided to the Semantic Precedent Retrieval Module, which uses vector representations and similarity-based retrieval to identify relevant court precedents.

The retrieved case laws and statutes are subsequently used by the LLM Synthesis Module to generate evidence-grounded legal summaries and insights. The system also evaluates grounded and ungrounded responses to identify unsupported information and fabricated precedents. Finally, the extracted findings, retrieved precedents, and generated summaries are presented through the

Streamlit interactive dashboard. This design provides a structured flow from legal document input to legal information extraction, precedent retrieval, evidence-supported generation, and final presentation.

3.3.1 Activity Diagram

The activity diagram represents the sequential workflow of the AI-Powered Legal Document Summarization and Case Precedent Finder. The process begins when the user uploads a legal document or provides legal text through the application. The document is parsed by the Document Parser, which converts the input into structured legal content. The processed content is then analyzed by Legal-BERT to extract important clauses, entities, rights, obligations, and liability-related information.

The extracted legal information is passed to the Semantic Precedent Retrieval Module, where Sentence-Transformers generate vector representations and similarity-based search identifies and ranks relevant court precedents. The retrieved case laws and statutes are then provided as supporting evidence to the Gemini-based synthesis component for generating a simplified legal summary. The system can evaluate the response as grounded or ungrounded to identify unsupported legal information.




Fig 3.2: End-to-End System Workflow of the Legal Document Analysis and Precedent Retrieval System

3.3.2 Data Flow Diagram

The Data Flow Diagram (DFD) illustrates how legal documents and related information are transformed into meaningful legal analysis outputs within the proposed AI-Powered Legal Document Summarization and Case Precedent Finder. The process begins with the user providing a legal document or legal text. The input is parsed and validated by the document processing module before being passed to the Legal-BERT Information Extraction Module.

The extracted legal information, including important clauses, entities, rights, obligations, and liability-related information, is then processed by the Semantic Precedent Retrieval Module. Sentence-Transformers convert the relevant legal content into vector representations, which are used to identify and rank relevant court precedents based on semantic similarity. The retrieved case precedents and statutes are then provided to the LLM Synthesis Module as supporting legal evidence. The generated summary and insights are designed to be grounded in the retrieved legal sources, while the system can evaluate grounded and ungrounded responses to identify unsupported information or fabricated precedents. Finally, the extracted findings.




Fig 3.3: Data Flow Diagram of the AI-Powered Legal Document Summarization and Case Precedent Finder

CHAPTER 4

PROJECT DESCRIPTION

4.1 METHODOLOGIES

The AI-Powered Legal Document Summarization and Case Precedent Finder employs an integrated Artificial Intelligence and Natural Language Processing methodology that combines Legal Document Processing, Legal-BERT-based Information Extraction, Semantic Precedent Retrieval, and Evidence-Grounded Summarization to improve legal document analysis and precedent discovery. The system first accepts legal documents in PDF and DOCX formats and processes them to extract meaningful legal content. The Legal-BERT module analyzes the processed text to identify important clauses, entities, rights, obligations, and liability-related information.

After document processing and extraction, the system applies semantic precedent retrieval to identify relevant court cases. Sentence-Transformers are used to convert legal content into meaningful vector representations, and similarity-based retrieval is applied to find and rank relevant precedents according to the context and meaning of the legal document. This approach allows the system to move beyond simple keyword-based searching and identify legally related cases based on semantic similarity.

The retrieved case laws and statutes are then used as supporting legal evidence for generating summaries and legal insights. The system uses Gemini for the synthesis stage, with the objective of producing evidence-grounded responses and reducing unsupported or fabricated legal information. The methodology also incorporates grounded-versus-ungrounded evaluation to identify fabricated precedents and unsupported statutes, thereby providing a mechanism for assessing the reliability of generated responses.

The methodology follows a modular architecture in which the document parser, Legal-BERT extraction module, retrieval layer, and LLM synthesis component perform separate but interconnected tasks. The final extracted findings, retrieved precedents, and generated summaries are presented through a Streamlit interactive dashboard.

4.2 MODULES

The AI-Powered Legal Document Summarization and Case Precedent Finder is organized into a set of interconnected modules, with each module responsible for a specific stage of the legal analysis workflow. The modular approach improves flexibility, simplifies maintenance, and allows individual components to be enhanced independently. Based on the proposed system architecture, the major modules are:

1. Document Processing Module
2. Legal-BERT Information Extraction Module
3. Semantic Precedent Retrieval Module
4. Evidence-Grounded Summarization Module
5. Interactive Dashboard Module

4.2.1 DOCUMENT PROCESSING MODULE

The Document Processing Module serves as the initial stage of the proposed system. It accepts legal documents in PDF and DOCX formats and converts their contents into structured legal text for further processing. The module is responsible for extracting relevant textual information while preparing the document for analysis by downstream components.

The processed legal content is passed to the Legal-BERT module for identifying important clauses, entities, rights, obligations, and liability-related information. This preprocessing stage provides consistent and structured input for subsequent legal information extraction and semantic retrieval. The document parser is implemented through the `doc_parser.py` component in the system architecture.

4.2.2 LEGAL-BERT INFORMATION EXTRACTION MODULE

The Legal-BERT Information Extraction Module forms an important component of the legal analysis pipeline. It uses Legal-BERT, a transformer-based language model adapted specifically for the legal domain, to understand the content and terminology of legal documents.

The module analyzes the processed legal text and extracts relevant information such as:

- Important legal clauses
- Legal entities
- Rights and obligations
- Liability-related information
- Relevant legal content

The extracted information is subsequently used by the precedent retrieval module to identify legally related cases. The project architecture identifies `legal_bert.py` as the component responsible for this stage.

4.1.3 SEMANTIC PRECEDENT RETRIEVAL MODULE

The Semantic Precedent Retrieval Module is responsible for finding relevant court cases based on the meaning and context of the legal document. Instead of depending only on exact keyword matching, the system converts legal content into meaningful vector representations using Sentence-Transformers.

The generated vectors are compared using cosine similarity to identify and rank relevant court precedents. The most semantically related case laws are then selected as supporting legal evidence for the subsequent summarization stage. This module is implemented through the `retrieval.py` component shown in the system architecture.

4.1.4 EVIDENCE-GROUNDED SUMMARIZATION MODULE

The Evidence-Grounded Summarization Module uses the retrieved legal precedents and statutes as supporting evidence for generating simplified legal summaries and insights. The project uses Gemini as the LLM synthesis component.

The main purpose of this module is to ensure that generated information is supported by retrieved legal evidence rather than relying only on the model's internal knowledge. This approach helps reduce unsupported claims and fabricated case

precedents. The project also includes a grounded-versus-ungrounded evaluation approach for identifying unsupported statutes and fabricated precedents.

4.1.5 INTERACTIVE DASHBOARD MODULE

The Interactive Dashboard Module provides the user interface for interacting with the proposed legal analysis system. The project uses Streamlit to present the results in an accessible and organized manner.

The dashboard is designed to provide users with:

- Extracted legal findings
- Important clauses and entities
- Retrieved case precedents
- Evidence supporting the retrieved information
- AI-generated summaries
- Legal analysis results

CHAPTER 5

IMPLEMENTATION AND RESULT DISCUSSION

5.1 IMPLEMENTATION

The implementation of the AI-Powered Legal Document Summarization and Case Precedent Finder is planned as the next major phase following the completion of the system design and architecture. The system will be implemented as a modular pipeline consisting of document processing, Legal-BERT-based information extraction, semantic precedent retrieval, and evidence-grounded summarization.

The implementation will begin with the development of the Document Processing Module, which accepts legal documents in PDF and DOCX formats and converts them into structured legal text. The extracted content will then be processed by the Legal-BERT Module to identify important clauses, entities, rights, obligations, and liability-related information. The project architecture identifies `doc_parser.py` and `legal_bert.py` as the primary components for these stages.

The next stage involves implementing the Semantic Precedent Retrieval Module using Sentence-Transformers. Legal content will be converted into vector representations and compared using similarity-based retrieval to identify and rank relevant court precedents. The retrieved precedents and statutes will then be passed to the Gemini-based synthesis component to generate evidence-grounded summaries.

The final stage will integrate the processing modules with the Streamlit interactive dashboard, allowing users to upload documents and view extracted findings, relevant precedents, and generated summaries. Each module will be developed and tested independently before being integrated into the complete workflow. The current project status identifies document parsing and evidence

retrieval as live components, while the Gemini-based grounded/ungrounded summary generation is currently marked as mocked.

5.1.1 Implementation of Data Pre-processing:

The Data Pre-processing Module acts as the initial processing stage of the proposed AI-Powered Legal Document Summarization and Case Precedent Finder. It prepares unstructured legal documents and extracted textual content before they are provided to the Legal-BERT and semantic retrieval components. Since legal documents may contain different formatting styles, lengthy clauses, special characters, and varied terminology, preprocessing is performed to obtain clean and consistent textual input.



Fig 5.1 Implementation UI Design

Legal Document Text Extraction and Sanitization:

Legal documents uploaded in PDF and DOCX formats are first parsed to extract their textual content. Unnecessary formatting elements and irrelevant characters are removed where required, while important legal terminology, clauses, entities, and contextual information are preserved. This ensures that the legal meaning of the original document is retained during subsequent processing. The project uses the doc_parser.py component for document parsing.

Textual pre-processing:

The extracted legal text is prepared for processing by the Legal-BERT model. The preprocessing stage organizes the document content into suitable textual units while preserving domain-specific legal terms and contextual relationships. This is important because Legal-BERT is subsequently used to identify important clauses, entities, rights, obligations, and liability-related information.




Fig 5.2 Text Input for Case Id Retrieval

Semantic Representation:

After preprocessing and legal information extraction, relevant legal content is converted into meaningful vector representations using Sentence-Transformers. These representations allow the system to compare the semantic meaning of the input legal content with available case precedents. Similarity-based retrieval is then used to identify and rank relevant precedents

Structured Data Preparation:

The processed information is organized into structured findings that can be passed between the extraction, retrieval, and summarization components. This structured representation supports the overall workflow from legal document processing to precedent retrieval and evidence-grounded generation.




Figure 5.3: Legal Document Data Pre-processing Pipeline

5.1.2 Implementation of Data Tokenization:

Once the legal document has been pre-processed, the textual content must be transformed into a representation that can be processed by the Legal-BERT model. The tokenization stage divides legal text into smaller token units and converts the resulting tokens into numerical representations that can be provided as model input.

- **Sub-Word Legal Text Segmentation:**

Legal text is divided into smaller sub-word tokens using the tokenizer associated with the Legal-BERT model. Sub-word tokenization enables the system to handle complex legal terminology, uncommon words, and domain-specific expressions even when complete words are not frequently encountered in the training vocabulary.

- **Sequence Standardization and Padding:**

Since transformer models process inputs in a standardized tensor format, tokenized legal text sequences are padded or truncated to the required context length. Attention masks are generated to distinguish actual tokens from padding tokens. This enables consistent batch processing of legal documents through the Legal-BERT model.

- **Numerical Token Representation:**

Each token is mapped to its corresponding numerical token ID according to the Legal-BERT vocabulary. These token IDs, together with the associated attention masks, form the model input used for subsequent legal text analysis and information extraction.

5.1.2 Implementation of Data Embedding and Feature Vectorization:

To enable the proposed AI-Powered Legal Document Summarization and Case Precedent Finder to understand the semantic meaning of legal text, the tokenized legal content is transformed into dense numerical vector representations. These embeddings capture relationships between legal terms and documents and provide the foundation for semantic precedent retrieval.

Dense Semantic Legal Embeddings:

The tokenized legal text is converted into dense vector representations using Sentence-Transformers. Instead of representing a legal document only as individual words or keyword counts, the embedding captures the contextual and semantic meaning of the legal content. Similar legal clauses and concepts therefore obtain similar vector representations, which supports more meaningful precedent retrieval.

Contextual Legal Feature Representation:

The Legal-BERT module processes tokenized legal text to identify meaningful legal features such as clauses, entities, rights, obligations, and liability-related information. These extracted features provide a structured representation of the legal document and can subsequently be used for legal analysis and precedent discovery.

Vector-Based Similarity Representation:

The generated embeddings represent legal content in a high-dimensional vector space. When a legal document or extracted finding is compared with stored case-law representations, similarity measures can be used to identify semantically

related cases. This enables the retrieval system to identify relevant precedents even when the wording of the input and the retrieved case is different.

Semantic Feature Integration:

The processed legal features and semantic embeddings are passed to the Semantic Precedent Retrieval Module. The retrieval component compares the vector representation of the input with available legal documents and ranks the most relevant precedents.




Figure 5.4: Legal Text Embedding and Feature Vectorization Architecture

EXPECTED RESULT:

The performance of the proposed AI-Powered Legal Document Summarization and Case Precedent Finder is expected to be evaluated through systematic testing of its major components, including legal information extraction, semantic precedent retrieval, and evidence-grounded summarization. The system will be assessed using a collection of legal documents and relevant case-law data to determine its effectiveness in processing and understanding legal content.

The Legal-BERT module is expected to accurately identify important legal information such as clauses, entities, rights, obligations, and liability-related content. The Semantic Precedent Retrieval Module is expected to identify and rank relevant court cases based on semantic similarity rather than relying only on exact keyword matching.

The generated summaries will be evaluated based on their groundedness and factual support, particularly by checking whether the generated legal statements are supported by retrieved statutes and case precedents. The project also includes evaluation of grounded versus ungrounded responses, including identification of unsupported statutes and fabricated precedents.

The expected outcome is a system capable of providing relevant legal findings, semantically related case precedents, and evidence-supported summaries through an interactive Streamlit dashboard. The overall evaluation is intended to demonstrate that the proposed framework can reduce manual effort in legal document analysis and support more structured and evidence-oriented legal research.

5.2.1 Quantitative Performance Benchmarks

Evaluation Metric	Baseline Model (Traditional FCFS)	MedLink Integrated Model	Operational Gain / Net Improvement
Wait-Time Error (MAE)	18.4 mins	3.1 mins	83.1% Reduction in Prediction Error
Overall Slot Utilization Rate	64.2%	91.8%	+27.6% Absolute Capacity Gain
Peak-Hour Wait-Time Reduction	48.6 mins	12.4 mins	74.4% Lower Patient Delay
System Disruption Recovery Time	> 120 mins	≤ 3 operational cycles	Near-Instantaneous Queue Re-convergence
API End-to-End Latency	1,240 ms	182 ms	85.3% Faster Response Velocity

Table 5.5: Performance Evaluation Matrix

5.2.2 Detailed Evaluation Analysis:




Fig 5.5: Wait-Time Prediction Error (MAE in Mins) Chart

Predictive Accuracy (Wait-Time Forecasting): The combined 1D-CNN + Bi-LSTM prediction architecture achieved a Mean Absolute Error (MAE) of 3.1 minutes, compared to 18.4 minutes observed in static block-scheduling systems. The initial 1D-CNN layers successfully isolated local delay spikes, while the Bi-LSTM layers captured long-term daily temporal trends.

Resource and Capacity Optimization: By implementing dynamic token generation and dynamic load balancing, overall clinic slot utilization rose to 91.8% (up from the baseline 64.2%). The platform effectively eliminated unabsorbed no-show gaps by auto-filling vacant slots with walk-in triage patients.

Peak Congestion Mitigation: During peak morning arrival surges (09:00 AM – 11:30 AM), average patient wait times were reduced by 74.4% (falling from

48.6 minutes down to 12.4 minutes). This performance improvement stems directly from the automated re-routing engine, which redirects non-critical patients when local clinic wait times exceed the 30-minute threshold.

5.2 DISCUSSION

The proposed AI-Powered Legal Document Summarization and Case Precedent Finder integrates domain-specific language understanding, semantic retrieval, and evidence-grounded generation to provide a structured approach to legal document analysis.

Effectiveness of Legal-BERT for Legal Information Extraction:

The Legal-BERT component is designed to improve the understanding of legal-domain terminology and document structure. Unlike general-purpose language models, Legal-BERT is specifically adapted to legal corpora and is used in the proposed system to identify important clauses, entities, rights, obligations, and liability-related information. This provides structured legal findings that can be used by subsequent modules for precedent retrieval and summarization.

Impact of Semantic Precedent Retrieval:

Traditional legal searches can depend heavily on keyword matching, which may fail when two legally related documents use different terminology. The proposed retrieval module addresses this by using Sentence-Transformers to generate vector representations and similarity-based retrieval to identify and rank relevant court precedents. This enables the system to consider the semantic relationship between the input legal document and available case-law information.

Role of Evidence-Grounded Generation:

The retrieved case laws and statutes provide supporting evidence for the LLM synthesis stage. The proposed approach is intended to reduce unsupported legal statements by grounding generated summaries in retrieved legal sources. The project also evaluates grounded and ungrounded responses, including unsupported statutes and fabricated precedents, to assess the reliability of generated legal

information.

Integration of the Legal Analysis Pipeline:

The main strength of the proposed approach is the integration of the individual components into a single workflow. Legal documents are first parsed, analyzed using Legal-BERT, processed into semantic representations, and used to retrieve relevant precedents before generating the final summary. The results are presented through the Streamlit dashboard, allowing extracted findings, precedents, and generated outputs to be accessed from a unified interface.

Current Implementation Consideration:

According to the project status, document parsing and evidence retrieval are live components, while the grounded/ungrounded Gemini summary generation is currently marked as mocked. Therefore, the final performance of the complete generation pipeline should be evaluated after the synthesis component is fully implemented.

CHAPTER 6

CONCLUSION AND FUTURE WORK

6.1 CONCLUSION

The proposed AI-Powered Legal Document Summarization and Case Precedent Finder presents an intelligent and integrated framework designed to improve the efficiency of legal document analysis and precedent discovery. The system combines legal document processing, Legal-BERT-based information extraction, semantic case precedent retrieval, and evidence-grounded summarization to address challenges such as time-consuming manual document analysis, difficulty in identifying relevant precedents, and unsupported information in AI-generated legal responses. The proposed architecture establishes a modular and scalable platform for processing legal documents and presenting meaningful legal findings through an interactive dashboard.

At this stage, the project has established its system requirements, architecture, data flow, activity workflow, and methodology for implementation. The selected techniques are designed to extract important legal clauses, entities, rights, obligations, and liability-related information using Legal-BERT, followed by semantic retrieval of relevant court precedents using Sentence-Transformers. The retrieved statutes and case laws are then intended to support the generation of evidence-grounded legal summaries.

The system also emphasizes evidence-based legal analysis, where retrieved legal sources are used to support generated outputs. The project incorporates grounded-versus-ungrounded evaluation to examine unsupported statutes and fabricated precedents. A Streamlit interactive dashboard provides the presentation layer for extracted findings, retrieved precedents, and generated summaries.

In summary, the proposed system provides a strong foundation for an AI-assisted legal research and document analysis platform. By combining domain-

specific language understanding, semantic precedent retrieval, and evidence-grounded generation, the system is expected to reduce manual effort in legal document analysis and make relevant legal information easier to identify and understand. Its modular architecture also provides flexibility for future improvements, including additional legal datasets, improved retrieval mechanisms, and enhanced evaluation of generated legal outputs.

6.2 FUTURE ENHANCEMENTS

As the next step, the AI-Powered Legal Document Summarization and Case Precedent Finder will progress into the implementation, integration, and testing stage, where the document processing, Legal-BERT extraction, semantic retrieval, and evidence-grounded summarization components will be further developed, integrated, validated, and optimized.

Future enhancements will include improving the Legal-BERT-based information extraction module to identify a wider range of legal entities, clauses, rights, obligations, and liability-related information from different types of legal documents. The semantic retrieval component can also be enhanced by expanding the legal case-law database and improving vector-based search and ranking to retrieve more relevant precedents.

Further improvements will involve strengthening the evidence-grounded generation and evaluation mechanisms to reduce unsupported legal statements and fabricated precedents. The system can be extended with additional legal datasets, improved retrieval models, and more comprehensive evaluation methods for measuring the relevance and reliability of generated summaries.

The Streamlit dashboard can also be enhanced to provide more interactive visualization of extracted legal findings, retrieved precedents, supporting

evidence, and generated summaries. Additional document formats and larger legal repositories can be incorporated to increase the system's applicability.

These advancements can transform the proposed framework into a more comprehensive, scalable, and reliable AI-assisted legal research and document analysis platform. Continuous evaluation and refinement of the individual modules can further improve system performance while maintaining the modular architecture established in the proposed design.

REFERENCES

- [1] I. Beltagy, M. E. Peters, and A. Cohan, “Longformer: The Long-Document Transformer,” *arXiv preprint arXiv:2004.05150*, 2020.
- [2] A. Bhattacharya, A. Ghosh, R. Ghosh, and S. Ghosh, “Deep Learning for Legal Text Summarization,” *Proc. International Conference on Artificial Intelligence and Law*, pp. 1–10, 2020.
- [3] N. A. Bommarito II and D. M. Katz, “A Mathematical Approach to the Study of the United States Code,” *Physical Review Letters*, vol. 107, no. 6, 2011.
- [4] M. Bommarito, D. Katz, and E. Zelner, “The Technology of Law: Legal Information Retrieval and Computational Legal Research,” *Artificial Intelligence and Law*, vol. 28, no. 1, pp. 1–20, 2020.
- [5] I. Chalkidis, A. Jana, D. Hartung, M. Bommarito, I. Androutsopoulos, D. Katz, and N. Aletras, “LexGLUE: A Benchmark Dataset for Legal Language Understanding in English,” *Proc. 60th Annual Meeting of the Association for Computational Linguistics*, pp. 4310–4330, May 2022.
- [6] I. Chalkidis, M. Fergadiotis, A. Savov, I. Androutsopoulos, and N. Aletras, “Paragraph-Level Rationale Extraction through Regularization: A Case Study on European Law,” *Proc. 16th Conference of the European Chapter of the Association for Computational Linguistics*, pp. 195–206, 2021.
- [7] I. Chalkidis, M. Fergadiotis, P. Malakasiotis, N. Aletras, and I. Androutsopoulos, “LEGAL-BERT: The Muppets straight out of Law School,” *Findings of the Association for Computational Linguistics: EMNLP 2020*, pp. 2898–2904, Nov. 2020.
- [8] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of Deep

Bidirectional Transformers for Language Understanding,” *Proc. NAACL-HLT*, vol. 1, pp. 4171–4186, Jun. 2019.

[9] M. Guo, Y. Zhang, and J. Zhao, “Legal Text Classification Using Pre-trained Language Models,” *Proc. International Conference on Artificial Intelligence and Law*, pp. 145–154, 2021.

[10] K. Guu, K. Lee, Z. Tung, P. Pasupat, and M. Chang, “Retrieval Augmented Language Model Pre-Training,” *Proc. 37th International Conference on Machine Learning*, vol. 119, pp. 3929–3938, 2020.

[11] K. Guu, K. Lee, Z. Tung, P. Pasupat, and M. Chang, “REALM: Retrieval-Augmented Language Model Pre-Training,” *Proc. 37th International Conference on Machine Learning*, vol. 119, pp. 3929–3938, 2020.

[12] M. Heddaya, K. MacMillan, H. Mei, C. Tan, and A. Malani, “CaseSumm: A Large-Scale Dataset for Long-Context Summarization from U.S. Supreme Court Opinions,” *Findings of the Association for Computational Linguistics: NAACL 2025*, pp. 1917–1942, Apr. 2025.

[13] J. Johnson, M. Douze, and H. Jégou, “Billion-Scale Similarity Search with GPUs,” *IEEE Transactions on Big Data*, vol. 7, no. 3, pp. 535–547, Sep. 2021.

[14] D. M. Katz, M. J. Bommarito II, J. Blackman, and J. D. D. Shah, “A General Approach for Predicting the Behavior of the Supreme Court of the United States,” *PLoS ONE*, vol. 12, no. 4, 2017.

[15] M. Lewis, Y. Liu, N. Goyal, M. Ghazvininejad, A. Mohamed, O. Levy, V. Stoyanov, and L. Zettlemoyer, “BART: Denoising Sequence-to-Sequence Pre-training for Natural Language Generation, Translation, and Comprehension,” *Proc. 58th Annual Meeting of*

the Association for Computational Linguistics, pp. 7871–7880, Jul. 2020.

[16] N. Reimers and I. Gurevych, “Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks,” *Proc. EMNLP-IJCNLP*, pp. 3982–3992, Nov. 2019.

[17] Y. Shao, J. Mao, and Y. Liu, “Legal Document Retrieval with Deep Neural Representation Learning,” *Proc. International Conference on Artificial Intelligence and Law*, pp. 125–134, 2021.

[18] Z. Shen, K. Lo, L. Yu, N. Dahlberg, M. Schlanger, and D. Downey, “Multi-LexSum: Real-World Summaries of Civil Rights Lawsuits at Multiple Granularities,” *arXiv preprint arXiv:2206.10883*, 2022.

[19] L. Zheng, N. Guha, B. R. Anderson, P. Henderson, and D. E. Ho, “When Does Pretraining Help? Assessing Self-Supervised Learning for Law and the CaseHOLD Dataset of 53,000+ Legal Holdings,” *Proc. 18th International Conference on Artificial Intelligence and Law (ICAIL)*, pp. 159–168, Jun. 2021.

[20] H. Zhong, Z. Guo, C. Tu, C. Xiao, Z. Liu, and M. Sun, “Legal Judgment Prediction via Topological Learning,” *Proc. 2018 Conference on Empirical Methods in Natural Language Processing*, pp. 3540–3549, Oct.–Nov. 2018.

LawXplorer_paper.docx



Indian Institute of Science

Document Details

Submission ID

trn:oid:::9619:56709298

Submission Date

Oct 5, 2026, 8:41 PM

Download Date

Oct 5, 2026, 8:41 PM

File Name

LawXplorer_paper.docx

File Size

1.1 MB

6 Pages

2,779 Words

17,212 Characters

8% Overall Similarity

The combined total of all matches, including overlapping sources, for each database.

Filtered from the Report

▸ Bibliography

Match Groups

- 

33 Not Cited or Quoted 8%

Matches with neither in-text citation nor quotation marks
- 

0 Missing Quotations 0%

Matches that are still very similar to source material
- 

0 Missing Citation 0%

Matches that have quotation marks, but no in-text citation
- 

0 Cited and Quoted 0%

Matches with in-text citation present, but no quotation marks

Top Sources

- 6%

 Internet sources
- 3%

 Publications
- 0%

 Submitted works

Match Groups

- 33 Not Cited or Quoted 8%**
 Matches with neither in-text citation nor quotation marks
- 0 Missing Quotations 0%**
 Matches that are still very similar to source material
- 0 Missing Citation 0%**
 Matches that have quotation marks, but no in-text citation
- 0 Cited and Quoted 0%**
 Matches with in-text citation present, but no quotation marks

Top Sources

- 6% Internet sources
- 3% Publications
- 0% Submitted works

Top Sources

The sources with the highest number of matches within the submission. Overlapping sources will not be displayed.

1	Internet	arxiv.org — How Does NLP Benefit Legal System: A Summary of Legal Artificial Intelligence	3%
2	Internet	arxiv.org — AutoLAW: Augmented Legal Reasoning through Legal Precedent Prediction	2%
3	Internet	arxiv.org — Embedding-based Recommender System for Job to Candidate Matching on Scale	2%
4	Internet	arxiv.org — A Survey on Legal Question Answering Systems	1%
5	Internet	arxiv.org — Similar Cases Recommendation using Legal Knowledge Graphs	1%
6	Internet	semanticscholar.org — Answer Book Valuation Using Semantic Similarity	1%
7	Internet	arxiv.org	1%
8	Internet	semanticscholar.org	1%
9	Internet	arxiv.org	1%
10	Internet	arxiv.org — Knowledge-Based Legal Document Assembly	1%
11	Internet	arxiv.org	1%
12	Publications	pubmed.ncbi.nlm.nih.gov	1%
13	Internet	arxiv.org	<1%
14	Internet	semanticscholar.org	<1%
15	Internet	arxiv.org	<1%
16	Internet	arxiv.org	<1%
17	Internet	arxiv.org	<1%

18	Publications	pubmed.ncbi.nlm.nih.gov	<1%
19	Internet	arxiv.org	<1%
20	Publications	pubmed.ncbi.nlm.nih.gov — A general structure for legal arguments about evidence using Bayesian networks.	<1%
21	Internet	arxiv.org	<1%

AI-POWERED LEGAL DOCUMENT SUMMARIZATION AND CASE PRECEDANT FINDER

Akiladevi R

Department of Computer Science
and Engineering
Rajalakshmi Engineering
College Thandalam, Chennai,
India
akiladevi.r@rajalakshmi.edu.in

Sanjay Kishan D

Department of Computer Science
and Engineering
Rajalakshmi Engineering College
Thandalam, Chennai, India
230701287@rajalakshmi.edu.in

Sanjeev Aditya P

Department of Computer Science
and Engineering
Rajalakshmi Engineering College
Thandalam, Chennai, India
230701292@rajalakshmi.edu.in

Sanjay R

Department of Computer Science
and Engineering
Rajalakshmi Engineering College
Thandalam, Chennai, India
230701291@rajalakshmi.edu.in

Sanjay G

Department of Computer Science
and Engineering
Rajalakshmi Engineering College
Thandalam, Chennai, India
230701286@rajalakshmi.edu.in

Abstract - LawXplorer is a novel AI-powered legal and compliance assistant that aims to transform the interaction of people, businesses and legal professionals with complex legal information. Traditional legal research and compliance processes are laborious, require lots of manual labour and require in-depth topic knowledge to decipher large statutes and agreements. LawXplorer helps to overcome these obstacles by deploying Natural Language Processing (NLP), Machine Learning (ML) and Generative AI that enable intelligent automation within legal processes. The system represents

a full-fledged digital legal advisor that receives legal questions from users in natural language, analyzes the questions and extracts relevant laws, and produces context-aware answers with references to actual laws.

Keywords: AI-driven legal aide, Digital legal evolution, Natural Language Understanding (NLU) Artificial Intelligence (AI) through Machine Learning (ML)

I. INTRODUCTION

The legal industry is experiencing a rapid digital transformation driven by the explosion of data, complex regulations and the need for quicker legal results.

Every year a flood of new cases, rules and compliance requirements impose an undue knowledge burden on the lawyers and institutions. Manual legal research and

compliance monitoring is time consuming and labor intensive, leading to delays, inconsistencies and human bias.

Recent advances in Artificial Intelligence (AI) and Natural Language Processing (NLP) have enabled the automation of legal interpretation, document review and compliance monitoring. Artificial intelligence algorithms are getting better at reading, understanding and summarizing legal documents and giving them contextually relevant insights. In this paper we present LawXplorer, an advanced legal assistant aimed at easing the interaction between people and organizations and complex legal systems. It leverages Natural Language Processing (NLP), Machine Learning (ML) and Generative AI to extract, evaluate and explain legal materials, making legal knowledge more accessible, efficient and transparent for all.

II. LITERATURE REVIEW

The influence of Artificial Intelligence on the legal arena has been substantial, fostering the advancement of systems that can decipher, scrutinize, and create legal insights. Multiple research endeavors and publications have delved into legal NLP models, the automation of compliance, and reasoning frameworks powered by AI.

Chalkidis et al. (2020) proposed Legal-BERT in their study [1], a transformer model that utilizes legal

corpora, hence boosting performance in tasks such as case classification and interpretation of statutes. Model [2], LexGLUE (2022), refined this notion by generating a benchmark dataset aimed at testing legal NLP systems across various subtasks, encompassing contract analysis and judgment prediction. The work by Zheng et al. (2021) presented CaseHOLD, a dataset specifically created for citation reasoning, facilitating AI's ability to predict case outcomes grounded in precedent reasoning.

Further examinations in references [4] and [5] highlighted RegTech and the automation of compliance, where AI solutions support the mapping of legislation, risk assessment, and conformity to regulations like GDPR and the EU AI Act. Analyses demonstrate that the utilization of Retrieval-Augmented Generation (RAG) increases the credibility of AI-generated replies by rooting them in actual legal literature.

Current commercial solutions like Westlaw Precision, Lexis+ AI, and Harvey AI reflect the rising implementation of AI in legal practice, but also accentuate persisting obstacles including financial implications, transparency, and jurisdictional alignment. In unison, these efforts accentuate the demand for a system that is elucidative, approachable, and compliant—an aspect that LawXplorer intends to address by integrating academic progress with real-world utility.

Beyond these foundational works, newer analyses emphasize the need for explainability and governance when dealing with legal artificial intelligence. Investigators such as Veale and Borgesius (2021) examined the requirement that AI-oriented legal systems conform to ethical and regulatory norms so as to eliminate prejudice and maintain accountability. Similarly, the EU AI Act (2024) has implemented compliance obligations that differ by risk, establishing new norms for transparency in high-risk domains like law and finance. These innovations show that building AI for legal contexts is a technological problem as well as a goal shaped by policy and trust. LawXplorer positions itself in this changing paradigm, aiming to provide insights that are clear, accountable and auditable, while keeping human control and taking advantage of the efficiencies offered by automation.

III. METHODOLOGIES USED

LawXplorer is designed with a systematic and modular AI-driven methodology for compliance monitoring, legal text comprehension and predictive assessment. The framework uses a retrieval-augmented generation (RAG) architecture that combines information retrieval and generative modeling to produce accurate, source-supported

answers. Legal documents, statutes and case law are first pre-processed using tokenization, named entity recognition (NER) and context embedding with Legal-BERT, a transformer model fine-tuned on the legal domain. This framework provides scalability, on-demand retraining and immediate reaction, which allow LawXplorer to deliver transparent and reliable AI results in complex legal settings.

A. NATURAL LANGUAGE PROCESSING (NLP)

LawXplorer is powered by Natural Language Processing that enables it to understand, analyze and generate legal terminology. NLP methods are used to segment text, identify entities and contextualize legal terminology. The system leverages transformer-based architectures, such as Legal-BERT, to extract clauses, case citations, and compliance terminology from large legal documents. This methodology not only collects data but also understands the importance of such data in relation to the relevant legal field.

The embeddings generated are stored in a vectorized format using FAISS indexing for fast retrieval and contextual linking. Combined, these components result in the NLP system of LawXplorer being a comprehension engine, not merely a text processor. It can read, reason and summarize legal documents with high accuracy.




Fig. 1. Conceptual diagram of NLP-based legal text processing pipeline

B. RETRIEVAL-AUGMENTED GENERATION (RAG)

The Retrieval-Augmented Generation framework unifies search-oriented retrieval with generative artificial intelligence. It first performs semantic similarity search to retrieve relevant legal documents. Then, the retrieved documents are combined with a generative model to generate contextually appropriate replies. This hybrid approach greatly reduces the rate of hallucinations and guarantees that any output is supported by solid sources (e.g., statutes or case law).

RAG is the “knowledge nucleus” of LawXplorer, to ensure that its responses are accurate and clear.




Fig. 2. Flow diagram of RAG-based answer generation

C. COMPLIANCE MONITORING SYSTEM

The system monitors compliance databases worldwide and regionally (like GDPR, EU AI Act) constantly. It uses rule-based mapping and machine learning classifiers to identify risk levels, detect change and alert users to changes that may impact their organization. Compliance can be monitored proactively, not reported on reactively, as the system is constantly ingesting regulatory updates in real-time.




Fig. 3. Conceptual Diagram of NLP-Based Text Processing Pipeline

D. PREDICTIVE LEGAL ANALYTICS

Such datasets are CaseHOLD and COLIEE which are used by machine learning methods to predict possible case outcomes. The model analyzes the trends in precedents, arguments and rulings to provide probability driven insights for legal practitioners and institutions. The predictive engine is intended to augment human reasoning, not replace it, by providing

data-driven foresight and retrospective analysis to aid in decision-making.

E. EXPLAINABLE AI FRAMEWORK

LawXplorer features a transparency layer that visually connects each AI insight to its legal reference. This ensures transparency and accountability that are the hallmarks of legal technology. The XAI layer enables the users to follow and trust the auditing standards through explaining the process of reasoning with the attention visualization and citation tracing.

F. DATA SOURCES AND TRAINING PROCESS

Model training is based on specialized legal datasets, such as LexGLUE, and publicly available collections of cases. The files contain legislation, court decisions and regulations from different jurisdictions and also have a variety of contexts. It uses multilingual data, which makes it suitable for local and international settings and allows it to grow in the future in different legal systems. LawXplorer's multilingual, multi-jurisdictional methodology not only improves accuracy but also establishes LawXplorer as a globally scalable legal AI system capable of comprehending and applying the law across a variety of linguistic, cultural, and regulatory environments.

IV. IMPLEMENTATION

LawXplorer collects and structures various legal information such as statutes, court cases or new regulations. It is trained and evaluated on public legal datasets such as LexGLUE, CaseHOLD and CUAD (Contract Understanding Atticus Dataset).

The datasets are cleaned and pre-processed thoroughly to get quality data for model training. Raw text is pre-processed to remove redundancies in parts, formatting inconsistencies, noise from citations. We used the spaCy library for tokenization and lemmatization. For better contextual understanding, domain specific stop words like “Section”, “Clause”, “Whereas” etc are removed.




Fig. 4 Architecture Diagram of proposed model

A. DATASET COLLECTION AND DATA PREPARATION

Every legal document has some meta data (case title, court name, jurisdiction, year) associated with it and is stored in a structured database. This enables the Retrieval-Augmented Generation (RAG) pipeline to find relevant documents in the context for query processing. TVision can consume real time JSON feeds from regulatory sources like EU AI Act, GDPR, Indian IT Rules etc for compliance monitoring and real time data updates.

Finally, all text data are converted into a vector format with the help of Legal-BERT embeddings, which are stored in an FAISS vector database for fast semantic retrieval.

B. FEATURE EXTRACTION

We then extract features to convert the unstructured legal text into a machine-understandable representation.

The system uses Legal-BERT and LexGLUE to build dense numerical embeddings for each clause, paragraph or section, that capture semantic and contextual meaning.

NER models are used to extract important entities such as parties, dates, obligations and risk terms for contract analysis. Similarly, in case law prediction, textual features including issue statements, parts of judgments and reasoning are extracted and classified into hierarchical categories.

At the same time, the compliance engine analyzes incoming regulatory feeds for structured features,

relating obligations, risk levels and compliance timelines.

The extracted features are all used as inputs to downstream modules for reasoning, prediction and visualization.

TRAINING AND EVALUATION

During LawXplorer's training phase, we focus on tuning the Legal-BERT model and improving the RAG pipeline.

We fine-tune the Legal-BERT encoder on the legal question-answer pairs from the LexGLUE and CaseHOLD datasets using the Hugging Face Transformers and PyTorch. The model is trained for 15 epochs with a learning rate of 2e-5 using AdamW optimizer and cross entropy loss function.

First, it pulls relevant information to augment the generation and make it more accurate. Thus the output is not a generic summary but is based on actual statutes and judgments.

The case prediction module predicts the probabilities of the outcome of the case based on the historical case data and supervised ML algorithms trained on this data, such as Random Forest and XGBoost. The model is evaluated based on F1-score, Precision and Recall and achieves an average accuracy of 89.4% over 150 benchmark queries.

Finally, the system generated outputs are validated by comparison with the authoritative legal sources for factual accuracy and reliability.




Fig. 5 Input Image I

Generated Result: Interpretation of Article 21 through LawXplorer




Fig. 6 Input Image II

Generated Result: Judicial Interpretations and Expansive Scope of Article 21 of the Indian Constitution

D. QUERY PROCESSING AND RESPONSE GENERATION

Once trained and fine-tuned, the core operational component of LawXplorer is the query processing pipeline. It does sufficiently to solve the problem of user interaction. It transforms natural language legal questions into powerful, evidence-driven answers. Then the user inputs a legal question into the online interface. The system uses Natural Language Processing (NLP) for initial detection of key entities, such as law provisions, article numbers, case citations etc. Legal-BERT then transforms the query into an embedding vector that captures the meaning and context of the query.

V. CONCLUSION

The AI-Driven Legal Document Summarization and Case Precedent Locator is a sophisticated and integrated system designed to improve the efficiency of legal document examination and precedent finding. To overcome the problems of manual document processing being time-consuming, the difficulty of finding relevant case precedents and the absence of evidences in AI-generated legal responses, the system combines legal document processing, Legal-BERT-based information extraction, semantic retrieval of case precedents and evidence-based summarization. The proposed framework is a modular and scalable system that processes legal documents and provides valuable legal insights using an interactive dashboard.

The project had defined system requirements, architecture, data flow, activity workflow and implementation strategy. Legal-BERT based techniques are employed to extract important legal clauses, entities, rights, obligations and liability related information which is then semantically

searched for relevant court precedents with Sentence-Transformers. The acquired case law and statutes are then used to create evidence based summaries of the law.

The system is evidence-based legal analysis, based on sourced legal materials to support its outputs. The initiative is assessing unfounded laws and fictitious precedents by comparing grounded and ungrounded judgments. The interactive Streamlit dashboard presents extracted findings, retrieved precedents and generated summaries.

Finally, the proposed system provides a good basis for an AI-powered legal research and document analysis platform. The system is designed to reduce the manual effort in analyzing legal documents by combining specialized language understanding, semantic extraction of precedents and evidence-based generation, enabling the identification and understanding of relevant legal information. In addition, the modular architecture provides flexibility for future improvements such as additional legal datasets, improved retrieval systems, and improved evaluation of the generated legal outputs.

The system intelligently uses NLP (Natural Language Processing), RAG (Retrieval Augmented Generation) and ML (Machine Learning) to understand user queries in natural language and generates contextually relevant responses referring to actual legal sources.

This is an initiative which takes a step toward democratizing data-centric legal technology by connecting crude legal data to valuable insight. The LawXplorer can be transformed into a complete digital companion to help legal practitioners, businesses and citizens understand, comply and feel confident in their legal choices, with further enhancements such as multi-lingual support, integration with official legal repositories and voice command capabilities.

REFERENCES

- [1] S. Bhosale, R. Chakraborty and S. K. Kopparapu, "A Novel Metric for Evaluating Audio Caption Similarity," in *2023 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2023, pp. 1–5.
- [2] Q. Fu, Y. Liu and Z. Xie, "EECS442 Final Project Report," University of Michigan, 2019.
- [3] C. Gao, Y. Dong, X. Yuan, Y. Han and H. Liu, "Infrared Image Captioning with Wearable Device," in *2023 IEEE International Conference on Robotics and Automation (ICRA)*, 2023, pp. 8187–8193.

[4] M. Hasegawa-Johnson, A. Black, L. Ondel, O. Scharenborg and F. Ciannella, "Image2Speech: Automatically Generating Audio Descriptions of Images," *The International Conference on Natural Language, Signal and Speech Processing*, Dec. 2017.

[5] Hrga and M. Ivašić-Kos, "Deep Image Captioning: An Overview," in *2019 42nd International Convention on Information and Communication Technology, Electronics and Microelectronics (MIPRO)*, 2019.

[6] K. Kandasamy, V. Padmanabhan, R. Priya Vallinayagam and S. K. Rajendran, "Survey on Chaos RNN—A Root Cause Analysis and Anomaly Detection," *AIP Conference Proceedings*, vol. 2790, no. 1, AIP Publishing, 2023.

[7] S. Mundargi and H. Mohanty, "Image Captioning Using Attention Mechanism with ResNet, VGG and Inception Models," *International Research Journal of Engineering and Technology (IRJET)*, vol. 7, no. 9, Sep. 2020.

[8] D. Naik and J. CD. Jaidhar, "Video Captioning Using Sentence Vector-Enabled Convolutional Framework with Short-Connected LSTM," *Multimedia Tools and Applications*, 2023, pp. 1–27.

[9] X. Ning, X. Zheng, Y. Yuan and X. Lu, "Audio Description from Image by Modal Translation Network," *Neurocomputing*, vol. 423, Jan. 2021.

[10] J. Y. Pan, H. J. Yang, P. Duygulu and C. Faloutsos, "Automatic Image Captioning," in *2004*

IEEE International Conference on Multimedia and Expo (ICME), 2004.

[11] R. Babu, K. Jayashree, K. B. Viswanathan and V. K., "An Efficient Spam Detector Model for Accurate Categorization of Spam Tweets Using Quantum Chaotic Optimization-Based Stacked Recurrent Network," *Nonlinear Dynamics*, vol. 111, no. 19, pp. 18523–18540, Oct. 2023, doi: 10.1007/s11071-023-08697-z.

[12] S. Rawale, M. Ghotkar, K. Sonavane, P. Surve, S. Khonde and D. Patil, "Image Captioning Generator System with Caption to Speech Conversion Mechanism," *International Research Journal of Modernization in Engineering Technology and Science*, vol. 3, no. 1, Jan. 2021.

[13] O. Vinyals, A. Toshev, S. Bengio and D. Erhan, "Show and Tell: A Neural Image Caption Generator," in *2015 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2015.

[14] H. Won, B. Kim, I.-Y. Kwak and C. Lim, "Using Various Pre-Trained Models for Audio Feature Extraction in Automated Audio Captioning," *Expert Systems with Applications*, 2023, Art. no. 120664.

[15] O. Yadav, A. Jadhav, A. H. Sunsara and I. Vohra, "Audio Caption Generation from Images Using Deep Learning," *International Journal of Trendy Research in Engineering and Technology*, vol. 5, no. 1, Feb. 2021. *Conference on Robotics and Automation (ICRA)* (pp. 8187-8193). IEEE.

LawXplorer_paper.docx



Indian Institute of Science

Document Details

Submission ID

trn:oid::961 9:56709298

Submission Date

Oct 05, 2026, 08:40 PM GMT+5:30

Download Date

Oct 05, 2026, 08:40 PM GMT+5:30

File Name

LawXplorer_paper.docx

File Size

1.1 MB

6 Pages

2,797 Words

20,011 Characters

0% detected as AI

The percentage indicates the combined amount of likely AI-generated text as well as likely AI-generated text that was also likely AI-paraphrased.

Caution: Review required.

It is essential to understand the limitations of AI detection before making decisions about a student's work. We encourage you to learn more about AI detection capabilities before using the tool.

DETECTION GROUPS



0 AI-generated only 0%

Likely AI-generated text from a large-language model.

Disclaimer

Our AI writing assessment is designed to help educators identify text that might be prepared by a generative AI tool. Our AI writing assessment may not always be accurate (it may misidentify writing that is likely AI generated as AI generated and AI paraphrased or likely AI generated and AI paraphrased writing as only AI generated) so it should not be used as the sole basis for adverse actions against a student. It takes further scrutiny and human judgment in conjunction with an organization's application of its specific academic policies to determine whether any academic misconduct has occurred.

AI-POWERED LEGAL DOCUMENT SUMMARIZATION AND CASE PRECEDANT FINDER

Akiladevi R

Department of Computer Science
and Engineering
Rajalakshmi Engineering College
Thandalam, Chennai, India
akiladevi.r@rajalakshmi.edu.in

Sanjay Kishan D

Department of Computer Science
and Engineering
Rajalakshmi Engineering College
Thandalam, Chennai, India
230701287@rajalakshmi.edu.in

Sanjeev Aditiya P

Department of Computer Science
and Engineering
Rajalakshmi Engineering College
Thandalam, Chennai, India
230701292@rajalakshmi.edu.in

Sanjay R

Department of Computer Science
and Engineering
Rajalakshmi Engineering College
Thandalam, Chennai, India
230701291@rajalakshmi.edu.in

Sanjay G

Department of Computer Science
and Engineering
Rajalakshmi Engineering College
Thandalam, Chennai, India
230701286@rajalakshmi.edu.in

Abstract - LawXplorer is a novel AI-powered legal and compliance assistant that aims to transform the interaction of people, businesses and legal professionals with complex legal information. Traditional legal research and compliance processes are laborious, require lots of manual labour and require in-depth topic knowledge to decipher large statutes and agreements. LawXplorer helps to overcome these obstacles by deploying Natural Language Processing (NLP), Machine Learning (ML) and Generative AI that enable intelligent automation within legal processes. The system represents a full-fledged digital legal advisor that receives legal questions from users in natural language, analyzes the questions and extracts relevant laws, and produces context-aware answers with references to actual laws.

Keywords: AI-driven legal aide, Digital legal evolution, Natural Language Understanding (NLU) Artificial Intelligence (AI) through Machine Learning (ML)

I. INTRODUCTION

The legal industry is experiencing a rapid digital transformation driven by the explosion of data, complex regulations and the need for quicker legal results.

Every year a flood of new cases, rules and compliance requirements impose an undue knowledge burden on the lawyers and institutions. Manual legal research and

compliance monitoring is time consuming and labor intensive, leading to delays, inconsistencies and human bias.

Recent advances in Artificial Intelligence (AI) and Natural Language Processing (NLP) have enabled the automation of legal interpretation, document review and compliance monitoring. Artificial intelligence algorithms are getting better at reading, understanding and summarizing legal documents and giving them contextually relevant insights. In this paper we present LawXplorer, an advanced legal assistant aimed at easing the interaction between people and organizations and complex legal systems. It leverages Natural Language Processing (NLP), Machine Learning (ML) and Generative AI to extract, evaluate and explain legal materials, making legal knowledge more accessible, efficient and transparent for all.

II. LITERATURE REVIEW

The influence of Artificial Intelligence on the legal arena has been substantial, fostering the advancement of systems that can decipher, scrutinize, and create legal insights. Multiple research endeavors and publications have delved into legal NLP models, the automation of compliance, and reasoning frameworks powered by AI.

Chalkidis et al. (2020) proposed Legal-BERT in their study [1], a transformer model that utilizes legal

corpora, hence boosting performance in tasks such as case classification and interpretation of statutes. Model [2], LexGLUE (2022), refined this notion by generating a benchmark dataset aimed at testing legal NLP systems across various subtasks, encompassing contract analysis and judgment prediction. The work by Zheng et al. (2021) presented CaseHOLD, a dataset specifically created for citation reasoning, facilitating AI's ability to predict case outcomes grounded in precedent reasoning.

Further examinations in references [4] and [5] highlighted RegTech and the automation of compliance, where AI solutions support the mapping of legislation, risk assessment, and conformity to regulations like GDPR and the EU AI Act. Analyses demonstrate that the utilization of Retrieval-Augmented Generation (RAG) increases the credibility of AI-generated replies by rooting them in actual legal literature.

Current commercial solutions like Westlaw Precision, Lexis+ AI, and Harvey AI reflect the rising implementation of AI in legal practice, but also accentuate persisting obstacles including financial implications, transparency, and jurisdictional alignment. In unison, these efforts accentuate the demand for a system that is elucidative, approachable, and compliant—an aspect that LawXplorer intends to address by integrating academic progress with real-world utility.

Beyond these foundational works, newer analyses emphasize the need for explainability and governance when dealing with legal artificial intelligence. Investigators such as Veale and Borgesius (2021) examined the requirement that AI-oriented legal systems conform to ethical and regulatory norms so as to eliminate prejudice and maintain accountability. Similarly, the EU AI Act (2024) has implemented compliance obligations that differ by risk, establishing new norms for transparency in high-risk domains like law and finance. These innovations show that building AI for legal contexts is a technological problem as well as a goal shaped by policy and trust. LawXplorer positions itself in this changing paradigm, aiming to provide insights that are clear, accountable and auditable, while keeping human control and taking advantage of the efficiencies offered by automation.

III. METHODOLOGIES USED

LawXplorer is designed with a systematic and modular AI-driven methodology for compliance monitoring, legal text comprehension and predictive assessment. The framework uses a retrieval-augmented generation (RAG) architecture that combines information retrieval and generative modeling to produce accurate, source-supported

answers. Legal documents, statutes and case law are first pre-processed using tokenization, named entity recognition (NER) and context embedding with Legal-BERT, a transformer model fine-tuned on the legal domain. This framework provides scalability, on-demand retraining and immediate reaction, which allow LawXplorer to deliver transparent and reliable AI results in complex legal settings.

A. NATURAL LANGUAGE PROCESSING (NLP)

LawXplorer is powered by Natural Language Processing that enables it to understand, analyze and generate legal terminology. NLP methods are used to segment text, identify entities and contextualize legal terminology. The system leverages transformer-based architectures, such as Legal-BERT, to extract clauses, case citations, and compliance terminology from large legal documents. This methodology not only collects data but also understands the importance of such data in relation to the relevant legal field.

The embeddings generated are stored in a vectorized format using FAISS indexing for fast retrieval and contextual linking. Combined, these components result in the NLP system of LawXplorer being a comprehension engine, not merely a text processor. It can read, reason and summarize legal documents with high accuracy.




Fig. 1. Conceptual diagram of NLP-based legal text processing pipeline

B. RETRIEVAL-AUGMENTED GENERATION (RAG)

The Retrieval-Augmented Generation framework unifies search-oriented retrieval with generative artificial intelligence. It first performs semantic similarity search to retrieve relevant legal documents. Then, the retrieved documents are combined with a generative model to generate contextually appropriate replies. This hybrid approach greatly reduces the rate of hallucinations and guarantees that any output is supported by solid sources (e.g., statutes or case law).

RAG is the “knowledge nucleus” of LawXplorer, to ensure that its responses are accurate and clear.




Fig. 2. Flow diagram of RAG-based answer generation

C. COMPLIANCE MONITORING SYSTEM

The system monitors compliance databases worldwide and regionally (like GDPR, EU AI Act) constantly. It uses rule-based mapping and machine learning classifiers to identify risk levels, detect change and alert users to changes that may impact their organization. Compliance can be monitored proactively, not reported on reactively, as the system is constantly ingesting regulatory updates in real-time.




Fig. 3. Conceptual Diagram of NLP-Based Text Processing Pipeline

D. PREDICTIVE LEGAL ANALYTICS

Such datasets are CaseHOLD and COLIEE which are used by machine learning methods to predict possible case outcomes. The model analyzes the trends in precedents, arguments and rulings to provide probability driven insights for legal practitioners and institutions. The predictive engine is intended to augment human reasoning, not replace it, by providing

data-driven foresight and retrospective analysis to aid in decision-making.

E. EXPLAINABLE AI FRAMEWORK

LawXplorer features a transparency layer that visually connects each AI insight to its legal reference. This ensures transparency and accountability that are the hallmarks of legal technology. The XAI layer enables the users to follow and trust the auditing standards through explaining the process of reasoning with the attention visualization and citation tracing.

F. DATA SOURCES AND TRAINING PROCESS

Model training is based on specialized legal datasets, such as LexGLUE, and publicly available collections of cases. The files contain legislation, court decisions and regulations from different jurisdictions and also have a variety of contexts. It uses multilingual data, which makes it suitable for local and international settings and allows it to grow in the future in different legal systems. LawXplorer's multilingual, multi-jurisdictional methodology not only improves accuracy but also establishes LawXplorer as a globally scalable legal AI system capable of comprehending and applying the law across a variety of linguistic, cultural, and regulatory environments.

IV. IMPLEMENTATION

LawXplorer collects and structures various legal information such as statutes, court cases or new regulations. It is trained and evaluated on public legal datasets such as LexGLUE, CaseHOLD and CUAD (Contract Understanding Atticus Dataset).

The datasets are cleaned and pre-processed thoroughly to get quality data for model training. Raw text is pre-processed to remove redundancies in parts, formatting inconsistencies, noise from citations. We used the spaCy library for tokenization and lemmatization. For better contextual understanding, domain specific stop words like “Section”, “Clause”, “Whereas” etc are removed.




Fig. 4 Architecture Diagram of proposed model

A. DATASET COLLECTION AND DATA PREPARATION

Every legal document has some meta data (case title, court name, jurisdiction, year) associated with it and is stored in a structured database. This enables the Retrieval-Augmented Generation (RAG) pipeline to find relevant documents in the context for query processing. TVision can consume real time JSON feeds from regulatory sources like EU AI Act, GDPR, Indian IT Rules etc for compliance monitoring and real time data updates. Finally, all text data are converted into a vector format with the help of Legal-BERT embeddings, which are stored in an FAISS vector database for fast semantic retrieval.

B. FEATURE EXTRACTION

We then extract features to convert the unstructured legal text into a machine-understandable representation.

The system uses Legal-BERT and LexGLUE to build dense numerical embeddings for each clause, paragraph or section, that capture semantic and contextual meaning.

NER models are used to extract important entities such as parties, dates, obligations and risk terms for contract analysis. Similarly, in case law prediction, textual features including issue statements, parts of judgments and reasoning are extracted and classified into hierarchical categories.

At the same time, the compliance engine analyzes incoming regulatory feeds for structured features,

relating obligations, risk levels and compliance timelines.

The extracted features are all used as inputs to downstream modules for reasoning, prediction and visualization.

TRAINING AND EVALUATION

During LawXplorer's training phase, we focus on tuning the Legal-BERT model and improving the RAG pipeline.

We fine-tune the Legal-BERT encoder on the legal question-answer pairs from the LexGLUE and CaseHOLD datasets using the Hugging Face Transformers and PyTorch. The model is trained for 15 epochs with a learning rate of 2e-5 using AdamW optimizer and cross entropy loss function.

First, it pulls relevant information to augment the generation and make it more accurate. Thus the output is not a generic summary but is based on actual statutes and judgments.

The case prediction module predicts the probabilities of the outcome of the case based on the historical case data and supervised ML algorithms trained on this data, such as Random Forest and XGBoost. The model is evaluated based on F1-score, Precision and Recall and achieves an average accuracy of 89.4% over 150 benchmark queries.

Finally, the system generated outputs are validated by comparison with the authoritative legal sources for factual accuracy and reliability.



Fig. 5 Input Image I

Generated Result: Interpretation of Article 21 through LawXplorer




Fig. 6 Input Image II

Generated Result: Judicial Interpretations and Expansive Scope of Article 21 of the Indian Constitution

D. QUERY PROCESSING AND RESPONSE GENERATION

Once trained and fine-tuned, the core operational component of LawXplorer is the query processing pipeline. It does sufficiently to solve the problem of user interaction. It transforms natural language legal questions into powerful, evidence-driven answers. Then the user inputs a legal question into the online interface. The system uses Natural Language Processing (NLP) for initial detection of key entities, such as law provisions, article numbers, case citations etc. Legal-BERT then transforms the query into an embedding vector that captures the meaning and context of the query.

V. CONCLUSION

The AI-Driven Legal Document Summarization and Case Precedent Locator is a sophisticated and integrated system designed to improve the efficiency of legal document examination and precedent finding. To overcome the problems of manual document processing being time-consuming, the difficulty of finding relevant case precedents and the absence of evidences in AI-generated legal responses, the system combines legal document processing, Legal-BERT-based information extraction, semantic retrieval of case precedents and evidence-based summarization. The proposed framework is a modular and scalable system that processes legal documents and provides valuable legal insights using an interactive dashboard.

The project had defined system requirements, architecture, data flow, activity workflow and implementation strategy. Legal-BERT based techniques are employed to extract important legal clauses, entities, rights, obligations and liability related information which is then semantically

searched for relevant court precedents with Sentence-Transformers. The acquired case law and statutes are then used to create evidence based summaries of the law.

The system is evidence-based legal analysis, based on sourced legal materials to support its outputs. The initiative is assessing unfounded laws and fictitious precedents by comparing grounded and ungrounded judgments. The interactive Streamlit dashboard presents extracted findings, retrieved precedents and generated summaries.

Finally, the proposed system provides a good basis for an AI-powered legal research and document analysis platform. The system is designed to reduce the manual effort in analyzing legal documents by combining specialized language understanding, semantic extraction of precedents and evidence-based generation, enabling the identification and understanding of relevant legal information. In addition, the modular architecture provides flexibility for future improvements such as additional legal datasets, improved retrieval systems, and improved evaluation of the generated legal outputs.

The system intelligently uses NLP (Natural Language Processing), RAG (Retrieval Augmented Generation) and ML (Machine Learning) to understand user queries in natural language and generates contextually relevant responses referring to actual legal sources.

This is an initiative which takes a step toward democratizing data-centric legal technology by connecting crude legal data to valuable insight. The LawXplorer can be transformed into a complete digital companion to help legal practitioners, businesses and citizens understand, comply and feel confident in their legal choices, with further enhancements such as multi-lingual support, integration with official legal repositories and voice command capabilities.

REFERENCES

- [1] S. Bhosale, R. Chakraborty and S. K. Kopparapu, "A Novel Metric for Evaluating Audio Caption Similarity," in *2023 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2023, pp. 1–5.
- [2] Q. Fu, Y. Liu and Z. Xie, "EECS442 Final Project Report," University of Michigan, 2019.
- [3] C. Gao, Y. Dong, X. Yuan, Y. Han and H. Liu, "Infrared Image Captioning with Wearable Device," in *2023 IEEE International Conference on Robotics and Automation (ICRA)*, 2023, pp. 8187–8193.

- [4] M. Hasegawa-Johnson, A. Black, L. Ondel, O. Scharenborg and F. Ciannella, "Image2Speech: Automatically Generating Audio Descriptions of Images," *The International Conference on Natural Language, Signal and Speech Processing*, Dec. 2017.
- [5] Hrga and M. Ivašić-Kos, "Deep Image Captioning: An Overview," in *2019 42nd International Convention on Information and Communication Technology, Electronics and Microelectronics (MIPRO)*, 2019.
- [6] K. Kandasamy, V. Padmanabhan, R. Priya Vallinayagam and S. K. Rajendran, "Survey on Chaos RNN-A Root Cause Analysis and Anomaly Detection," *AIP Conference Proceedings*, vol. 2790, no. 1, AIP Publishing, 2023.
- [7] S. Mundargi and H. Mohanty, "Image Captioning Using Attention Mechanism with ResNet, VGG and Inception Models," *International Research Journal of Engineering and Technology (IRJET)*, vol. 7, no. 9, Sep. 2020.
- [8] D. Naik and J. CD. Jaidhar, "Video Captioning Using Sentence Vector-Enabled Convolutional Framework with Short-Connected LSTM," *Multimedia Tools and Applications*, 2023, pp. 1–27.
- [9] X. Ning, X. Zheng, Y. Yuan and X. Lu, "Audio Description from Image by Modal Translation Network," *Neurocomputing*, vol. 423, Jan. 2021.
- [10] J. Y. Pan, H. J. Yang, P. Duygulu and C. Faloutsos, "Automatic Image Captioning," in *2004 IEEE International Conference on Multimedia and Expo (ICME)*, 2004.
- [11] R. Babu, K. Jayashree, K. B. Viswanathan and V. K., "An Efficient Spam Detector Model for Accurate Categorization of Spam Tweets Using Quantum Chaotic Optimization-Based Stacked Recurrent Network," *Nonlinear Dynamics*, vol. 111, no. 19, pp. 18523–18540, Oct. 2023, doi: 10.1007/s11071-023-08697-z.
- [12] S. Rawale, M. Ghotkar, K. Sonavane, P. Surve, S. Khonde and D. Patil, "Image Captioning Generator System with Caption to Speech Conversion Mechanism," *International Research Journal of Modernization in Engineering Technology and Science*, vol. 3, no. 1, Jan. 2021.
- [13] O. Vinyals, A. Toshev, S. Bengio and D. Erhan, "Show and Tell: A Neural Image Caption Generator," in *2015 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2015.
- [14] H. Won, B. Kim, I.-Y. Kwak and C. Lim, "Using Various Pre-Trained Models for Audio Feature Extraction in Automated Audio Captioning," *Expert Systems with Applications*, 2023, Art. no. 120664.
- [15] O. Yadav, A. Jadhav, A. H. Sunsara and I. Vohra, "Audio Caption Generation from Images Using Deep Learning," *International Journal of Trendy Research in Engineering and Technology*, vol. 5, no. 1, Feb. 2021. *Conference on Robotics and Automation (ICRA)* (pp. 8187-8193). IEEE.